

# Fast Differential Dynamic Programming for Time-Optimal Trajectory Planning

Research Portfolio

Miaomiao Dai

<https://github.com/dmmsjtu-umich/time-opt-ilqr>

## 1 Overview

Time-optimal trajectory planning—generating motions that complete a task in the minimum possible time—is a fundamental requirement for agile robotic systems. From autonomous drone racing to emergency collision avoidance in self-driving cars, the ability to jointly optimize the control sequence and the total maneuver duration  $T$  is critical for pushing physical limits.

While Differential Dynamic Programming (DDP) and its variant, the iterative Linear Quadratic Regulator (iLQR), have become standard tools for high-dimensional trajectory optimization, they typically assume a fixed planning horizon. Extending these methods to time-optimal control introduces a discrete-continuous optimization challenge: the solver must determine the optimal integer horizon  $T^*$  alongside the continuous control inputs.

## 2 Problem Formulation

We consider the discrete-time optimal control problem with variable horizon:

$$\begin{aligned} \min_{U_T, T} \quad & J = \phi(x_T) + \sum_{k=0}^{T-1} \ell(x_k, u_k) + wT \\ \text{s.t.} \quad & x_{k+1} = f(x_k, u_k), \quad k = 0, \dots, T-1 \\ & x_0 = \bar{x}_0, \quad T \in \{1, 2, \dots, N\} \end{aligned} \tag{1}$$

where the decision variables include both the control sequence  $U_T = \{u_0, \dots, u_{T-1}\}$  and the finish time  $T$ . The term  $wT$  penalizes the planning horizon, encouraging time-optimal behaviour of the system.

## 3 The Challenge

**A fundamental bottleneck lies in the structure of the standard Riccati recursion.** In the LQR backward pass, the Value Function  $V_k$  is computed recursively starting from a terminal cost anchored at the final time step  $T$  (i.e.,  $P_T = Q_T$ ). Consequently, changing the horizon from  $T$  to  $T+1$  shifts the boundary condition, invalidating the entire sequence of previously computed Cost-to-Go matrices. This structural dependency prevents the reuse of historical computations across different horizons, forcing the solver to restart the backward pass from scratch for each candidate  $T$ , resulting in a prohibitive  $\mathcal{O}(N^2)$  complexity.

## 4 Our Approach: Time-Varying Propagator

To address the loss of reusability in time-varying systems, we shift from *reusing values* to *reusing mappings*.

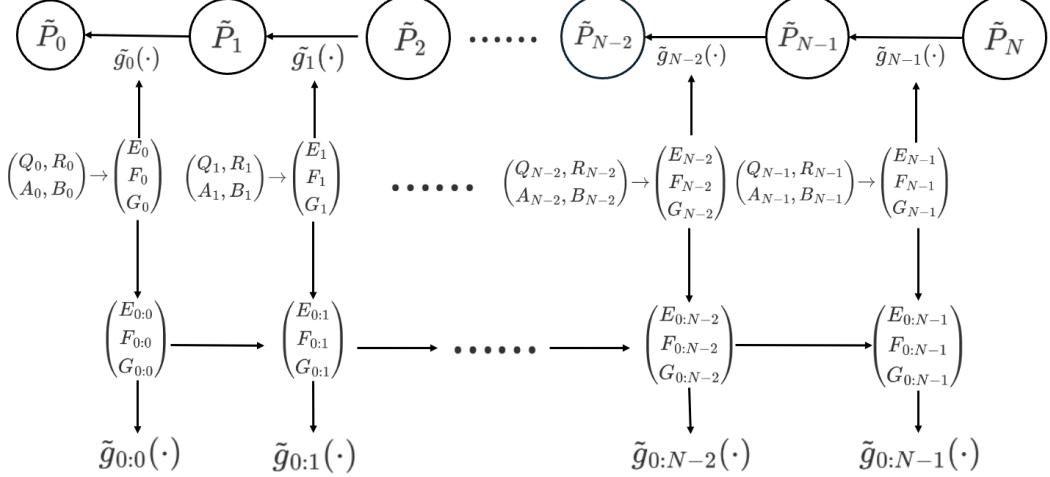


Figure 1: **Time-varying propagator.** When  $g_k$  varies with  $k$ , we switch to inverse form where each stage is an LFT  $\tilde{g}_k$ . The composed map  $\tilde{g}_{0:k}$  remains an LFT with prefix parameters  $(E_{0:k}, F_{0:k}, G_{0:k})$ . This enables cheap horizon queries by reusing the composed mapping  $\tilde{g}_{0:k}$  instead of reusing  $\tilde{P}_k$  values.

Our key idea (Fig. 1) is to rewrite the map  $g_k$  as a new linear fractional transformation (LFT) form  $\tilde{g}_{0:k}$ , and some of the matrices that help compute  $\tilde{g}_{0:k}$  can be reused. As a result, these matrices only need to be computed once for all possible horizons  $k = 1, 2, \dots, N$ , as opposed to be repetitively computed for each possible horizon, which thus saves computational effort.

### 4.1 Linear Fractional Transformation Form

Let  $\tilde{P}_k := P_k^{-1}$  denote the inverse matrix of  $P_k$ . Let notation  $\tilde{g}_{0:k} = \tilde{g}_0 \circ \dots \circ \tilde{g}_k$  denote a *composed map* that composes the maps  $g_0, g_1, \dots, g_k$  sequentially.

**Theorem 1** (LFT Form). *There exist matrices  $(E_{0:k}, F_{0:k}, G_{0:k})$ ,  $k = 0, 1, 2, \dots, N$  such that*

$$\tilde{g}_{0:k}(\tilde{P}) = E_{0:k} - F_{0:k}(\tilde{P} + G_{0:k})^{-1}F_{0:k}^\top, \quad (2)$$

where the prefix parameters obey the recursion (for  $k \geq 1$ ):

$$\begin{aligned} W_k &= (E_k + G_{0:k-1})^{-1}, \\ E_{0:k} &= E_{0:k-1} - F_{0:k-1}W_kF_{0:k-1}^\top, \\ F_{0:k} &= F_{0:k-1}W_kF_k, \\ G_{0:k} &= G_k - F_k^\top W_kF_k, \end{aligned} \quad (3)$$

with base case  $E_{0:0} = E_0$ ,  $F_{0:0} = F_0$ ,  $G_{0:0} = G_0$ , and stage parameters:

$$E_k = Q_k^{-1}, \quad F_k = Q_k^{-1}A_k^\top, \quad G_k = A_kQ_k^{-1}A_k^\top + B_kR_k^{-1}B_k^\top. \quad (4)$$

## 4.2 Efficient Query for All Arrival Times

Given the prefix triple at  $t - 1$ , the initial inverse for any candidate arrival time  $t$  is:

$$\tilde{P}_0^{(t)} = E_{0:t-1} - F_{0:t-1}(\tilde{P}_t + G_{0:t-1})^{-1}F_{0:t-1}^\top. \quad (5)$$

Then  $P_0^{(t)} = (\tilde{P}_0^{(t)})^{-1}$  and the cost for horizon  $t$  is:

$$J_t = \frac{1}{2}x_0^\top P_0^{(t)}x_0 + wt. \quad (6)$$

Thus all  $\{J_t\}_{t=1}^N$  are obtained from a single forward prefix build plus  $N$  terminal updates.

**Complexity.** The propagator has the same order as a single LQR backward sweep,  $\mathcal{O}(Nn^3)$ . Brute forcing all horizons by re-solving Riccati is  $\mathcal{O}(N^2n^3)$ .

## 4.3 Augmented State Formulation for iLQR

To extend the propagator to nonlinear iLQR, we introduce an **Augmented State Formulation** that absorbs the time-varying affine linearization terms into a homogeneous coordinate system:

$$z_k = \begin{bmatrix} \delta x_k \\ 1 \end{bmatrix}. \quad (7)$$

The augmented system matrices become:

$$A_k^{\text{aug}} = \begin{bmatrix} A_k - B_k \ell_{uu,k}^{-1} \ell_{ux,k} & -B_k \ell_{uu,k}^{-1} \ell_{u,k} \\ 0 & 1 \end{bmatrix}, \quad B_k^{\text{aug}} = \begin{bmatrix} B_k \\ 0 \end{bmatrix}. \quad (8)$$

This unifies the treatment of linear and nonlinear problems, allowing the propagator to compute the *exact* LQR cost for all horizons in a single  $\mathcal{O}(N)$  pass.

## 5 Main Contributions

1. **Propagator-based Horizon Selection:** We develop an LFT-based solver that enables the reuse of backward pass computations, reducing the complexity of horizon selection from  $\mathcal{O}(N^2n^3)$  to  $\mathcal{O}(Nn^3)$ .
2. **Augmented State Formulation:** We propose a state augmentation technique that embeds affine linearization terms into a homogeneous coordinate system, extending the efficient propagator method to general nonlinear iLQR problems.
3. **Performance and Robustness:** We validate our algorithm on four benchmark systems, including a 12-DOF Quadrotor. Experimental results show that our method achieves speedups of up to **43×** compared to brute-force search while guaranteeing global optimality with respect to the linearized model.

## 6 Results

We validate our proposed Propagator-based iLQR on four benchmark systems: Double Integrator, Segway Balance, Cartpole Swing-Up, and a 12-DOF Quadrotor. We compare three horizon-selection strategies:

- **Ours (Propagator):** The proposed method using the Augmented State Propagator.
- **Baseline-1 (Bruteforce):** Evaluating all horizons via standard backward Riccati sweeps. This serves as the ground truth but has  $\mathcal{O}(N^2)$  complexity.
- **Baseline-2 (OnePass):** The current state-of-the-art approximate method, which estimates costs for neighboring horizons by reusing the value function from a single nominal backward pass. While efficient, it is only an *approximation* that can introduce significant errors for time-varying systems.

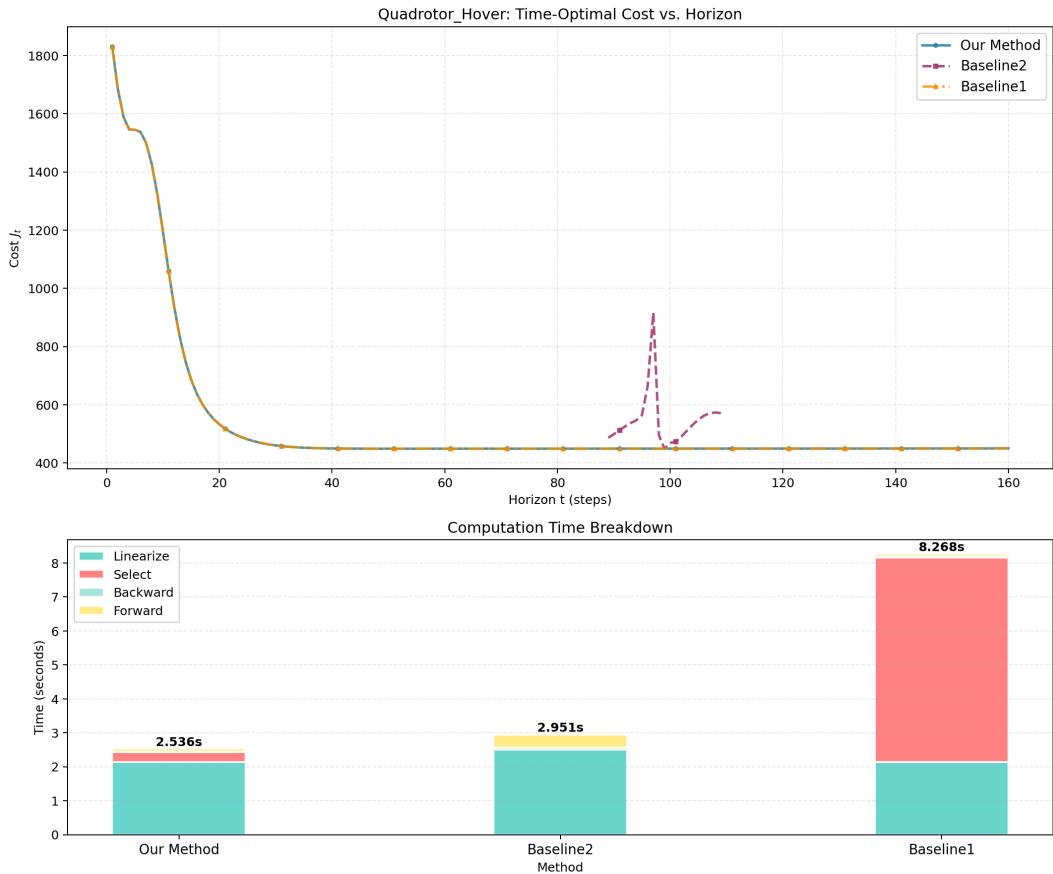


Figure 2: **Case Study on Quadrotor Hover.** (Top) Comparison of cost landscapes ( $J_t$ ) computed by different methods. (Bottom) Breakdown of total runtime into linearization, selection, backward, and forward phases.

**Cost Landscape (Top Panel).** The top panel of Fig. 2 compares the cost curves ( $J_t$  vs. horizon  $t$ ). The OnePass method (Purple), as the current SOTA, approximates the cost landscape by

projecting the value function from a single nominal horizon. However, because it is merely an *estimation* method, it suffers from severe distortions when the system dynamics vary significantly over time (e.g., the artifact spike near  $t = 82$ ). Consequently, OnePass converges to a wrong local minimum, failing to find the true optimal horizon.

In contrast, our Propagator curve (Blue) overlaps *perfectly* with the Bruteforce ground truth (Yellow). This confirms that our augmented formulation correctly captures the *exact* time-varying LQR cost—not an approximation—allowing the solver to locate the true global optimum.

**Runtime Breakdown (Bottom Panel).** The bottom panel decomposes the runtime to reveal the source of efficiency. The Bruteforce method (Baseline-1) computes the exact cost but at prohibitive expense—the massive “Select” phase (Red) represents the  $\mathcal{O}(N^2)$  cost of repeated Riccati sweeps, taking **8.21s** total.

The OnePass method (Baseline-2) is faster (**2.88s**) but, as shown above, produces incorrect results due to its approximate nature.

Our Propagator method achieves the best of both worlds: it computes the *exact* cost like Bruteforce while running even faster than the approximate OnePass (**2.46s**). By compressing the horizon evaluation into an  $\mathcal{O}(N)$  LFT propagation, we effectively eliminate the selection bottleneck while maintaining rigorous optimality.